

SRT411 Take Home Assignment

Michael Nguyen

March 19, 2016

Defining the Problem

Introduction

On March 18th, I let Wireshark run starting from 9:00am until 3:00pm, while instructing my younger brother that he may do whatever he liked on my computer. Some possible scenarios that could come out of this are captures of packets going to various different servers around the world, and a plentitude of possible protocols used. By the time I arrived home, the total packet capture was up to roughly 6.6 million packets and when exported added up to a whopping 1GB file. Knowing the limitation of knitting PDF, if i were to display all lines, by either viewing or printing this in one go would be very resource extensive on individuals trying to open this PDF. My goal is to try to figure out what happened during the time I left Wireshark on, until the time I arrived home and present this information in a way that shows that I uncovered some of this information.

Constructing Questions

Firstly after setting my working directory and importing the .csv file, i took a summary of the data in order to see what i was working with. Below is the code that was used: >data=read.csv("mar18")

```
>info<-summary(data)
>write.csv(info,"info.csv")
```

```
info=read.csv("info.csv")
print(info)
```

```
##      X          X.....No.      X.....Time      X.....Source
## 1 NA Min.      :      1  Min.      :      0  192.168.1.2      :2647794
## 2 NA 1st Qu.:1663104  1st Qu.: 1336  13.107.4.50      :2084780
## 3 NA Median :3326208  Median : 2126  69.28.210.33      : 292683
## 4 NA Mean   :3326208  Mean   : 5733  149.56.0.38       : 224188
## 5 NA 3rd Qu.:4989312  3rd Qu.: 8131  198.71.80.215     : 199342
## 6 NA Max.   :6652415  Max.   :20854  162.254.199.80    :180900
## 7 NA          <NA>          <NA> (Other)      :1022728
##      X.....Destination      X.....Protocol      X....Length
## 1 192.168.1.2      :3918079  TCP          :3484267  Min.      : 42.0
## 2 13.107.4.50      :1058182  GVSP         :2555242  1st Qu.:  78.0
## 3 149.56.0.38      : 224566  UDP          : 424644  Median : 194.0
## 4 198.71.80.215    : 199629  ARP          :  76420  Mean   : 615.5
## 5 162.254.199.80   :179217  HTTP         : 48429  3rd Qu.:1506.0
## 6 162.254.199.90   :146645  CLASSIC-STUN: 23381  Max.    :11594.0
## 7 (Other)         : 926097  (Other)      : 40032          <NA>
##      X.....Info
## 1 61233 > 28015  Len=11      : 105648
## 2 57138 > 27101  Len=11      :  83732
## 3 28015 > 61233  Len=52      :  72915
## 4 28015 > 61233  Len=11      :  22813
```

```
## 5 [TCP segment of a reassembled PDU]: 15804
## 6 27101 > 57138 Len=11 : 14215
## 7 (Other) :6337288
```

*#Since it would be too resource extensive
#to import the dataset into this PDF, workarounds had to be made.*

In order for us to even look at the summary in this PDF, I had to write the output from looking at summary function to a *csv* or another readable format by *R* and read that instead. This process would be very important in order to construct this document with easy accessibility.

Now, that we have a glance at the information at hand we can begin to develop questions. Some simple ones that could be examined are:

1. How many unique Destination IPs are there?
2. What Protocols were used?
3. Can we look at the frequency of packets throughout the day?
4. Is there any applications or conversations we can look at or identify?

And furthermore we question which way we can plot or visualize this data:

5. Can this data be represented on a map?
6. What constraints do we need to take into account for graphing?

I will attempt to answer these questions if-possible in visual representations, and explore other options where available.

Assessing Available Data

Dataset Information

Looking back at the summary of the dataset in the first part, we can see that there are **7 different variables**. These are as listed below: * No. (The packet number) * Time (ms) * Source IP * Destination IP * Protocol * Packet Length * Information Field of packet

Apart from that we can use our other functions such as summary and mutate to see counts and frequency or create new columns using the current existing ones. It is most likely

Other Sources

Geo-location

In order to solve one of the questions posted in the first section, such as determining geolocation, I will be using this function I found here , and the employment of the function below with a simple example:

```
library(rjson)
freegeoip <- function(ip, format = ifelse(length(ip)==1,'list','dataframe'))
{
  if (1 == length(ip))
  {
```

```

# a single IP address
require(rjson)
url <- paste(c("http://freegeoip.net/json/", ip), collapse='')
ret <- fromJSON(readLines(url, warn=FALSE))
if (format == 'dataframe')
  ret <- data.frame(t(unlist(ret)))
return(ret)
} else {
  ret <- data.frame()
  for (i in 1:length(ip))
  {
    r <- freegeoip(ip[i], format="dataframe")
    ret <- rbind(ret, r)
  }
  return(ret)
}
}

```

#Using the first non-private IP found in my dataset

```
freegeoip("13.107.4.50")
```

```

## $ip
## [1] "13.107.4.50"
##
## $country_code
## [1] "US"
##
## $country_name
## [1] "United States"
##
## $region_code
## [1] ""
##
## $region_name
## [1] ""
##
## $city
## [1] ""
##
## $zip_code
## [1] ""
##
## $time_zone
## [1] ""
##
## $latitude
## [1] 38
##
## $longitude
## [1] -97
##
## $metro_code

```

```
## [1] 0
```

With this we can figure out Longitude and Latitude of an *IP*, but we will need to play with our data a bit so we can feed this function more than one value at a time.

Processing Information

Now we proceed with information screening to develop results. There will several parts within this section documenting different methods or techniques used to filter the information from the dataset.

Destinations

First I want to look at the amount of unique destination addresses within the dataset. Using the code below:

```
udest <- unique(data$Destination)
```

```
length(udest$x)
```

```
## [1] 6286
```

```
#We find that there are 6286 unique results for the Destination Address  
#Will eventually need to see which ones are relevant
```

Again since I there is no dataset present in this PDF, normally by just taking the `length(udest)` from result of using the unique function on the **data\$Destination** .

Protocols Usage

To see a general list of which protocols were used we would also want to do a unique or summary of the protocols used.

```
proto<-summary(data$Protocol)
```

```
#Display of Protocols and Count and Frequency  
#We add the frequency and round it for easier workability.  
proto <-mutate(proto,freq = round(Count/66524.15,digits=3))  
proto
```

##	Protocol	Count	freq
## 1	ARP	76420	1.149
## 2	BROWSER	74	0.001
## 3	CAT-TP	1	0.000
## 4	CLASSIC-STUN	23381	0.351
## 5	DCERPC	1	0.000
## 6	DCP-PFT	9	0.000
## 7	DHCP	13	0.000
## 8	DHCPv6	12	0.000
## 9	DIS	2	0.000

## 10	DNS	2588	0.039
## 11	FTP	56	0.001
## 12	FTP-DATA	4	0.000
## 13	GVSP	2555242	38.411
## 14	HTTP	48429	0.728
## 15	HTTP/XML	588	0.009
## 16	ICMP	55	0.001
## 17	ICMPv6	4097	0.062
## 18	IGMPv3	334	0.005
## 19	IPv6	90	0.001
## 20	KNET	2	0.000
## 21	LANMAN	36	0.001
## 22	LLC	2	0.000
## 23	LLDP	22	0.000
## 24	LLMNR	920	0.014
## 25	MDNS	86	0.001
## 26	MP4	1	0.000
## 27	NBNS	253	0.004
## 28	NBSS	36	0.001
## 29	OCSP	7	0.000
## 30	Pathport	5	0.000
## 31	PNRP	1026	0.015
## 32	QUAKE	2	0.000
## 33	QUAKE2	1	0.000
## 34	QUIC	5274	0.079
## 35	RTCP	2002	0.030
## 36	SIGCOMP	2	0.000
## 37	SMB	117	0.002
## 38	SMB2	27	0.000
## 39	SNMP	323	0.005
## 40	SSDP	5056	0.076
## 41	SSL	3	0.000
## 42	TCP	3484267	52.376
## 43	TLSv1	451	0.007
## 44	TLSv1.2	16444	0.247
## 45	UAUDP	2	0.000
## 46	UDP	424644	6.383
## 47	XID	8	0.000

*#Note: it is possible to achieve similar data using verisr's getenum
#although format will be off.*

There are 47 different Protocols found within the capture. Although the information regarding frequency was rounded to 3 digits past the decimal, we still get a good picture of the overall results. The Protocol **GVSP** which is neither one of the two standard transport protocols **TCP** and **UDP** had a surprising amount of packets. Upon looking this up, it would be appropriate because **GVSP** is used for video streaming or VOIP using **UDP** packets.

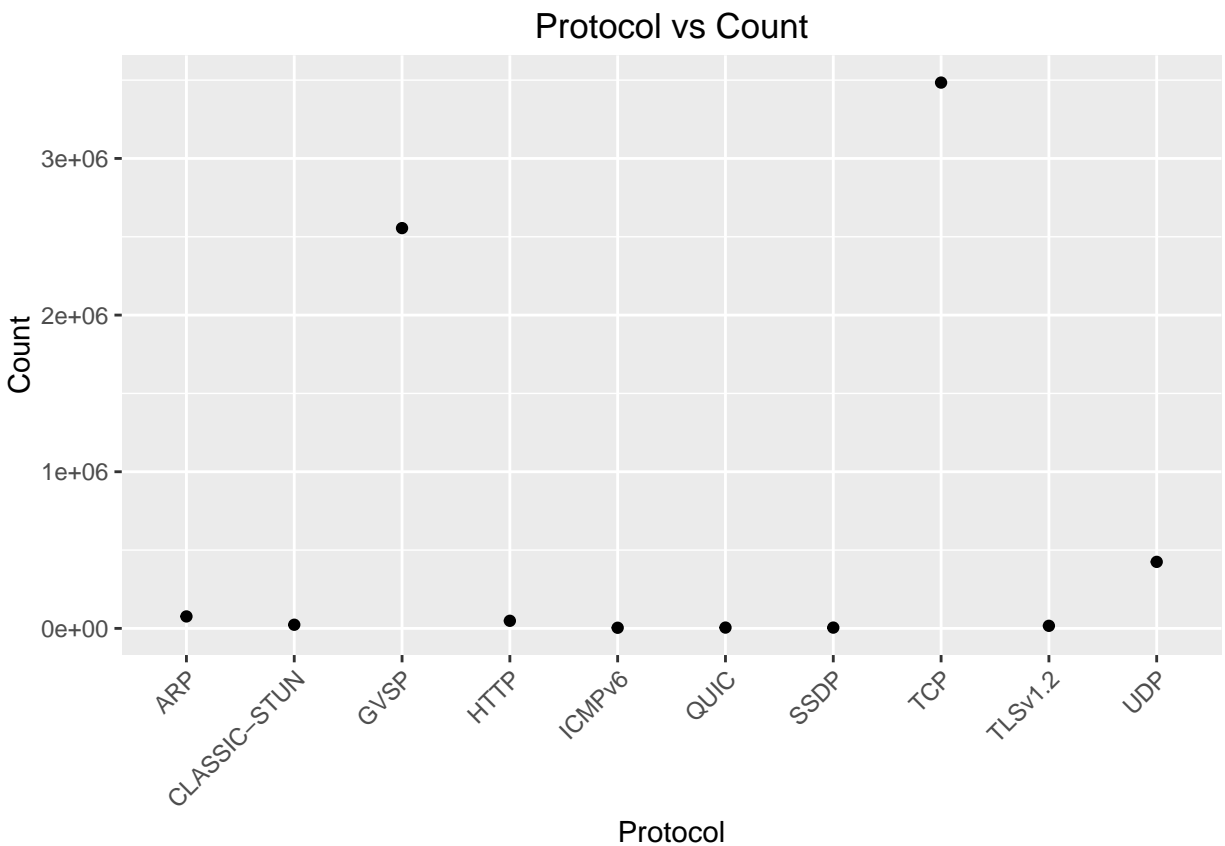
Visualizing & Viewing Transformation

Graphs

In this section we develop our initial graphics for the data that has been processed attempting to separate fields by use of graphics for easier readability and also ensure that these graphs do not overutilize resources within the system. We will also try to answer some of the questions posed in the first section.

A

```
#Graph using freq  
#Seperate top 10 Protocols  
proto1<-top_n(proto,10,Count)  
qplot(Protocol,Count,data=proto1)+theme(axis.text.x = element_text(angle = 45, hjust = 1))+ggtitle("Pro
```




That was just an overview of which protocols were most often used during the session, now we begin to look at the stand out protocol for me which was **GVSP**

```
gd <-select(data,Time,Protocol)  
gd <-filter(gd,Protocol=="GVSP")
```

B

```
g1<- ggplot(gd,aes(Time,Length))
g1<- g1+geom_smooth()
g1 <-g1+labs(title="Average GVSP Packet Length Over Time",xlabs="Time (ms)")
print(g1)
```



With approximately 38% or 2.3million packets, it would very intensive on any system to plot this information and convey it in a PDF. This graph plots the average length over time of GSVP packets, and is a rough approximation of how the application was sending and receiving packets. If this were plotted in either **geom_dotplot** or **geom_jitter** it would nearly impossible to load.

C

```
by_time <- group_by(data,Protocol,Destination)
sumtime<- summarise (by_time)
```

Now we will comb through the specific IPs used and the associated protocol that was the destination of the communication. By performing a **group_by** function we can separate the data by the fields we want. In this case I want to find out all the **HTTP** Destination IPs.

```
onlyhttp<-filter(sumtime,Protocol=="HTTP")
head(onlyhttp)
```

```
## Protocol Destination
## 1 HTTP 104.16.26.216
## 2 HTTP 104.193.82.53
## 3 HTTP 104.20.45.97
## 4 HTTP 104.20.46.97
## 5 HTTP 104.23.129.82
## 6 HTTP 104.238.216.36
```

Now we can use our **freegeoip** function to tack on a Longitude and Latitude to each IP, after making some modifications to the data frame.

```
dest1 <- select(data,Protocol,Destination)
dest1 <- filter(dest1,Protocol=="HTTP")
dest2 <- unique(dest1$Destination)
final<-freegeoip(dest2)
world<-filter(final,latitude!=0 & longitude!=0)
```

Getting a sample of the cleaned IPs, that now have geo-location

```
head(world)
```

```
## ip country_code country_name region_code region_name
## 1 206.248.168.136 CA Canada
## 2 72.21.91.8 US United States CA California
## 3 198.41.214.187 US United States CA California
## 4 23.76.117.127 US United States MA Massachusetts
## 5 206.248.168.177 CA Canada
## 6 184.86.40.154 US United States MA Massachusetts
## city zip_code time_zone latitude longitude metro_code
## 1 43.6425 -79.3873 0
## 2 Santa Monica 90405 America/Los_Angeles 34.0119 -118.4682 803
## 3 San Francisco 94107 America/Los_Angeles 37.7697 -122.3933 807
## 4 Cambridge 02142 America/New_York 42.3626 -71.0843 506
## 5 43.6425 -79.3873 0
## 6 Cambridge 02142 America/New_York 42.3626 -71.0843 506
```

```
library(rworldmap)
```

```
## Loading required package: sp
```

```
## ### Welcome to rworldmap ###
```

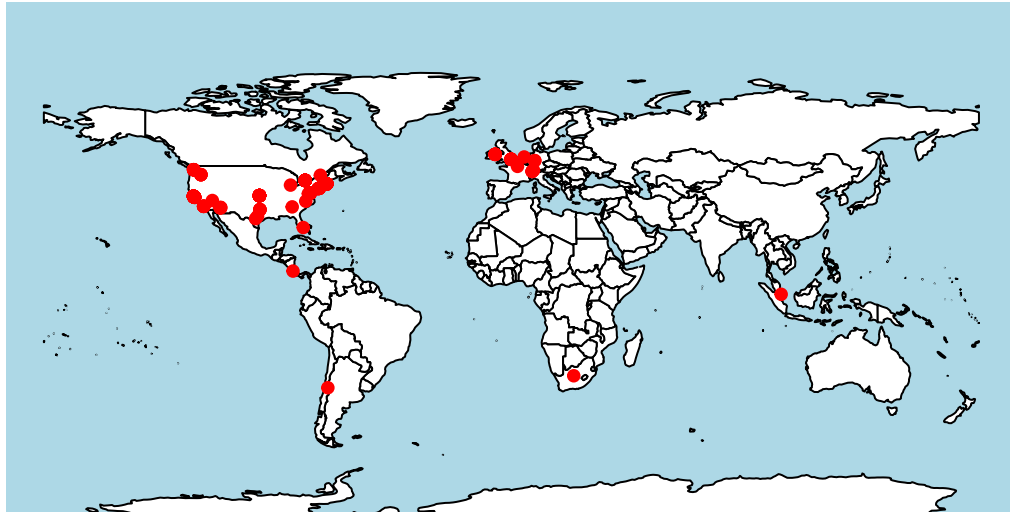
```
## For a short introduction type : vignette('rworldmap')
```

```
map <- getMap(resolution = "HIGH")
```

```
## Warning in getMap(resolution = "HIGH"): resolution should be set to one of :coarse low less islands
## setting to coarse as default
```

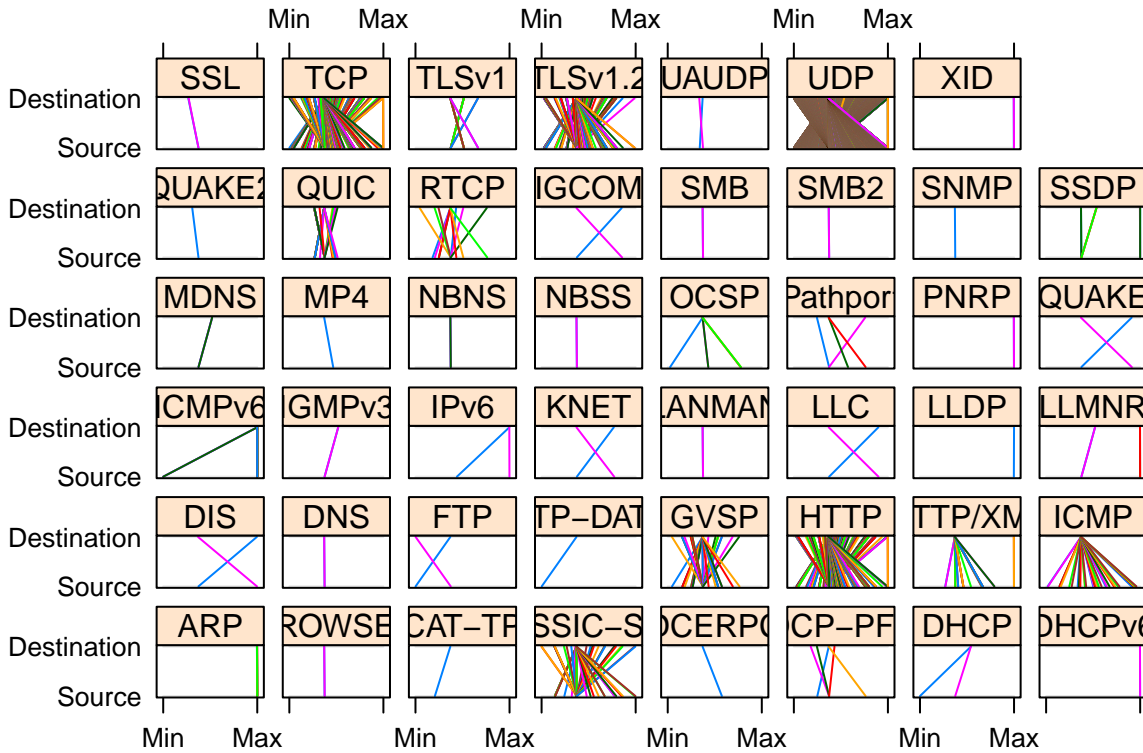


```
plot(map, bg="lightblue", col="white", ylim=c(-60, 90), mar=c(0,0,0,0))
points(world$longitude, world$latitude, col="red", cex=0.8, pch=19)
```



This plot of the Destination IPs and their relative location to the world shows where the **HTTP** packets travelled to. The majority reside in either the east or west coast of North America, however there were connections made to 5 continents in total (minus Antarctica, and Australia). We get a good idea of where web servers are located. ###D Attempting to show the stream of data depicting the relationship of Source and Destination IP grouped by Protocol

```
parallelplot(~extratime[1:2] | extratime$Protocol)
```




The column after the pipe adds the *facet* feature to the graph and allows us to graph each relationship in regards to its protocol. It is very clear and vivid in the graphs of **DHCP** and **OCSP**, where you can see that all the various source IPs are connected to the same destination IP.

E

Finally, in the last demonstration, we take a look at the packet length in comparison to the protocol used.

```
ss<-ggplot(vers,aes(Length,Protocol))
ss<- ss+geom_jitter(aes(color = Length))
ss<-ss+geom_smooth()
ss<- ss+theme(text=element_text(family="Times", size=6))
ss <-ss+labs(title="Protocol vs Packet Length")
print(ss)
```

```
## Warning: Computation failed in `stat_smooth()`:
## x has insufficient unique values to support 10 knots: reduce k.
```



Although, there are flags for error when using `geom_smooth`, with the addition of it, the graph dots from the different rows are not as tightly knit together, making it easier to distinguish the lengths of each packet within a protocol.

Conclusion and Final Thoughts

Apart from graphing the information, I tried very hard to keep this PDF as minimal resource-intensive as possible. I knew that using certain plots would drastically slow down the rendering of this PDF when opened. Although my dataset was quite large, I believe this method of reading smaller tables and files (usually hidden within R markdown) made it so the file was much quicker to access. Compiling with knitr itself was quite quick with times of under 30seconds. I will try to upload this, and view it on a different computer to check my hypothesis.

I believe i was able to answer the questions that were made in the first section; However I found it quite difficult to map out frequency for all protocols throughout the day. With a bit more work, I can improve on visual portions of the graphs but the constraint that is the dataset itself was ever-present within the work. Normally loading the dataset using `read.csv()` would take roughly 5~10mins and graphing anything related to dots would usually throw out an error. Although this is a first look for myself and big data, I think the data represented covered some of the topics and also things I found interesting.

The lack of a duration column within this dataset has got me thinking. With it and with such a large dataset we can use the duration to figure out on average how fast it takes for a protocol to function.

Notes

This `.rmd` file will most likely be up on my GitHub It can found at the following link